

Package ‘mvMonitoring’

February 2, 2017

Type Package

Title Multi-State Adaptive Dynamic Principle Components Analysis for
Multivariate Process Monitoring

Version 0.1.0

Author Gabriel Odom, Ben Barnard, and Karen Kazor

Maintainer Gabriel Odom <gabriel_odom@baylor.edu>

Description Use multi-state splitting to apply Adaptive-Dynamic PCA to data
generated from a continuous-time multivariate industrial or natural process.
Employ PCA-based dimension reduction to extract linear combinations of
relevant features, reducing computational burdens.

License file LICENSE

Depends R (>= 2.10)

Imports BMS, lazyeval, zoo, xts, utils

Encoding UTF-8

LazyData true

RoxygenNote 5.0.1

Suggests knitr,
rmarkdown

VignetteBuilder knitr

R topics documented:

fault1A_xts	2
fault1B_xts	3
fault2A_xts	3
fault2B_xts	4
fault3A_xts	5
fault3B_xts	6
faultDetect	6
faultFilter	7
mspMonitor	9
mspTrain	10
mspWarning	11
normal_switch_xts	13
pca	13

processMonitor	14
rotate3D	15
rotateScale3D	16
threshold	17
Index	18

 fault1A_xts

Process Data under a System Shift Fault

Description

Three-feature, three-state process data including observations under normal operating conditions and observations after a positive shift for each feature in the system.

Usage

fault1A_xts

Format

An xts data matrix with 10080 rows, corresponding to one week worth of data recorded at a 1-minute interval. The columns under normal conditions are defined in the help file for normal_switch_xts. The fault is a system shock to each of the three features by 2. The fault starts at row 8500, and the four columns under the fault state are defined here:

state the state indicator for the multivariate system, with three levels

x $x(t) = t + 2 + \text{error}$

y $y(t) = t^2 - 3t + 2 + \text{error}$

z $z(t) = -t^3 + 3t^2 + 2 + \text{error}$

where t is a 10080-entry vector of autocorrelated and non-stationary hidden process realizations. The states alternate each hour and are defined as follows:

State1 As presented

State2 Rotated by (yaw = 0, pitch = 90, roll = 30) and scaled by (1 * x, 0.5 * y, 2 * z).

State3 Rotated by (yaw = 90, pitch = 0, roll = -30) and scaled by (0.25 * x, 0.1 * y, 0.75 * z).

See the vignette for more details.

Source

Simluated in R.

 fault1B_xts

Process Data under a Feature Shift Fault

Description

Three-feature, three-state process data including observations under normal operating conditions and observations after a positive shift in values for one feature.

Usage

```
fault1B_xts
```

Format

An xts data matrix with 10080 rows, corresponding to one week worth of data recorded at a 1-minute interval. The columns under normal conditions are defined in the help file for normal_switch_xts. The fault is a system shock to the "x" feature only by 2. The fault starts at row 8500, and the four columns under the fault state are defined here:

state the state indicator for the multivariate system, with three levels

x $x(t) = t + 2 + \text{error}$

y $y(t) = t^2 - 3t + \text{error}$

z $z(t) = -t^3 + 3t^2 + \text{error}$

where t is a 10080-entry vector of autocorrelated and non-stationary hidden process realizations. The states alternate each hour and are defined as follows:

State1 As presented

State2 Rotated by (yaw = 0, pitch = 90, roll = 30) and scaled by $(1 * x, 0.5 * y, 2 * z)$.

State3 Rotated by (yaw = 90, pitch = 0, roll = -30) and scaled by $(0.25 * x, 0.1 * y, 0.75 * z)$.

See the vignette for more details.

Source

Simulated in R.

 fault2A_xts

Process Data under a System Drift Fault

Description

Three-feature, three-state process data including observations under normal operating conditions and observations after a positive drift in values for each feature in the system.

Usage

```
fault2A_xts
```

Format

An xts data matrix with 10080 rows, corresponding to one week worth of data recorded at a 1-minute interval. The columns under normal conditions are defined in the help file for normal_switch_xts. The fault is a drift on each feature by $s / 10^3$, where s is the observation index. The fault starts at row 8500, and the four columns under the fault state are defined here:

state the state indicator for the multivariate system, with three levels

x $x(t) = t + \text{drift} + \text{error}$

y $y(t) = t^2 - 3t + \text{drift} + \text{error}$

z $z(t) = -t^3 + 3t^2 + \text{drift} + \text{error}$

where t is a 10080-entry vector of autocorrelated and non-stationary hidden process realizations. The states alternate each hour and are defined as follows:

State1 As presented

State2 Rotated by (yaw = 0, pitch = 90, roll = 30) and scaled by $(1 * x, 0.5 * y, 2 * z)$.

State3 Rotated by (yaw = 90, pitch = 0, roll = -30) and scaled by $(0.25 * x, 0.1 * y, 0.75 * z)$.

See the vignette for more details.

Source

Simulated in R.

fault2B_xts

Process Data under a Feature Drift Fault

Description

Three-feature, three-state process data including observations under normal operating conditions and observations after a positive drift in values for two features.

Usage

fault2B_xts

Format

An xts data matrix with 10080 rows, corresponding to one week worth of data recorded at a 1-minute interval. The columns under normal conditions are defined in the help file for normal_switch_xts. The fault is a drift on the "y" and "z" features by $s / 10^3$, where s is the observation index. The fault starts at row 8500, and the four columns under the fault state are defined here:

state the state indicator for the multivariate system, with three levels

x $x(t) = t + \text{error}$

y $y(t) = t^2 - 3t + \text{drift} + \text{error}$

z $z(t) = -t^3 + 3t^2 + \text{drift} + \text{error}$

where t is a 10080-entry vector of autocorrelated and non-stationary hidden process realizations. The states alternate each hour and are defined as follows:

State1 As presented

State2 Rotated by (yaw = 0, pitch = 90, roll = 30) and scaled by (1 * x, 0.5 * y, 2 * z).

State3 Rotated by (yaw = 90, pitch = 0, roll = -30) and scaled by (0.25 * x, 0.1 * y, 0.75 * z).

See the vignette for more details.

Source

Simluated in R.

fault3A_xts

Process Data under a System Signal Amplification

Description

Three-feature, three-state process data including observations under normal operating conditions and observations after an amplification of the underlying process for each feature in the system.

Usage

fault3A_xts

Format

An xts data matrix with 10080 rows, corresponding to one week worth of data recorded at a 1-minute interval. The columns under normal conditions are defined in the help file for normal_switch_xts. The fault is a signal amplification in the underlying determining t vector. The fault starts at row 8500, and the four columns under the fault state are defined here:

state the state indicator for the multivariate system, with three levels

x $x(t) = t_* + \text{error}$

y $y(t) = (t_*)^2 - 3t + \text{error}$

z $z(t) = -(t_*)^3 + 3(t_*)^2 + \text{error}$

where $t_* = 3 * t * (10080 - s) / (2 * 10080)$, where s is the observation index, and t is a 10080-entry vector of autocorrelated and non-stationary hidden process realizations. The states alternate each hour and are defined as follows:

State1 As presented

State2 Rotated by (yaw = 0, pitch = 90, roll = 30) and scaled by (1 * x, 0.5 * y, 2 * z).

State3 Rotated by (yaw = 90, pitch = 0, roll = -30) and scaled by (0.25 * x, 0.1 * y, 0.75 * z).

See the vignette for more details.

Source

Simluated in R.

 fault3B_xts

Process Data under a Feature Signal Dampening

Description

Three-feature, three-state process data including observations under normal operating conditions and observations after a dampening of one of the underlying features in the system.

Usage

```
fault3B_xts
```

Format

An xts data matrix with 10080 rows, corresponding to one week worth of data recorded at a 1-minute interval. The columns under normal conditions are defined in the help file for normal_switch_xts. The fault is a signal dampening in the underlying determining t vector for the "z" feature. The fault starts at row 8500, and the four columns under the fault state are defined here:

state the state indicator for the multivariate system, with three levels

x $x(t) = t + \text{error}$

y $y(t) = t^2 - 3t + \text{error}$

z $z(t) = -(t_*)^3 + 3(t_*)^2 + \text{error}$

where $t_* = \log|t|$, and t is a 10080-entry vector of autocorrelated and non-stationary hidden process realizations. The states alternate each hour and are defined as follows:

State1 As presented

State2 Rotated by (yaw = 0, pitch = 90, roll = 30) and scaled by (1 * x, 0.5 * y, 2 * z).

State3 Rotated by (yaw = 90, pitch = 0, roll = -30) and scaled by (0.25 * x, 0.1 * y, 0.75 * z).

See the vignette for more details.

Source

Simulated in R.

 faultDetect

Process Fault Detection

Description

Detect if a single multivariate observation is beyond normal parameters.

Usage

```
faultDetect(threshold_object, observation, ...)
```

Arguments

threshold_object	An object of classes "threshold" and "pca" returned by the internal threshold() function.
observation	A single row of an xts data matrix to compare against the thresholds
...	Lazy dots for additional internal arguments

Details

This function takes in a threshold object returned by the threshold() function and a single observation as a matrix or xts row. The function then returns a row vector of the SPE test statistics, a logical indicator marking if this statistic is beyond the threshold, the Hotelling's T2 statistic, and an indicator if this statistic is beyond the threshold. Observations with monitoring statistics beyond the calculated threshold are marked with a 1, while within-parameter observations are marked with a 0. These threshold values are passed from the threshold() function through this function via a returned threshold object. This object will be used in higher function calls.

This internal function is called by faultFilter().

Value

A named 1 * 4 matrix of the SPE statistic value ("SPE"), SPE fault indicator ("SPE_Flag"), T2 statistic value ("T2"), and T2 fault indicator for the single row observation passed to this function ("T2_Flag").

Examples

```
data("normal_switch_xts")
scaledData <- scale(normal_switch_xts[, -1])
pca_obj <- pca(scaledData, var.amnt = 0.9)
thresh_obj <- threshold(pca_object = pca_obj, alpha = 0.05)

# Check a single observation. We see this observation is within the normal
# operating parameters at alpha = 0.05.
faultDetect(threshold_object = thresh_obj, observation = scaledData[1,])
# According to the Squared Prediction Error statistic, this observation is
# outside the range of "normal" operation at the 0.05 level.
faultDetect(threshold_object = thresh_obj, observation = scaledData[20,])

# We can also check an entire data matrix:
detect_ls <- lapply(1:nrow(scaledData), function(i){
  faultDetect(threshold_object = thresh_obj, scaledData[i,])
})
do.call(rbind, detect_ls)
```

Description

Flag and filter out observations beyond normal parameters, then return the observations within normal operating conditions.

Usage

```
faultFilter(trainData, testData, updateFreq, faultsToTriggerAlarm = 3, ...)
```

Arguments

<code>trainData</code>	An xts data matrix of initial training observations
<code>testData</code>	The data not included in the training data set
<code>updateFreq</code>	How many observations from the test data matrix will be returned to update the training data matrix and move it forward?
<code>faultsToTriggerAlarm</code>	Specifies how many sequential faults will cause an alarm to trigger. Defaults to 3.
<code>...</code>	Lazy dots for additional internal arguments

Details

This function is essentially a wrapper function to call and organize the output from these other internal functions: `faultDetect()`, `threshold()`, and `pca()`. It is applied over a rolling window, with observation width equal to `updateFreq`, of the larger full data matrix via the `processMonitor()` function, wherein the testing and training data sets move forward in time across the entire data matrix.

This internal function is called by `processMonitor()`.

Value

A list of class "fault_ls" with the following: `faultObj` = an xts flagging matrix with the same number of rows as "testData". This flag matrix has the following five columns: the SPE test statistic for each observation in "testData", an SPE indicator recording 0 if the test statistic is less than or equal to the critical value passed through from the threshold object, a T2 test statistic, a similar T2 indicator, and a final column indicating if there have been three flags in a row for either the SPE or T2 monitoring statistics or both. `nonAlarmedTestObs` = an xts matrix of the first `updateFreq` number of rows of the training data which were not alarmed. `trainSpecs` = the threshold object returned by the internal `threshold()` function. See this function's help file for more details.

Examples

```
data("normal_switch_xts")
# Select the data under state 1
data <- normal_switch_xts[normal_switch_xts[,1] == 1]

# Split the data into testing and training data sets
nTrainObs <- floor(0.2 * nrow(data))
nUpdate <- floor(0.5 * nTrainObs)
trainObs <- data[1:nTrainObs, -1]
testObs <- data[(nTrainObs + 1):nrow(data), -1]

faultFilter(trainData = trainObs,
            testData = testObs,
            updateFreq = nUpdate)
```


mspMonitor

*Real-Time Process Monitoring Function***Description**

Monitor and flag (if necessary) incoming multivariate process observations.

Usage

```
mspMonitor(observations, labelVector, trainingSummary, ...)
```

Arguments

observations	an $n * p$ xts matrix. For real-time monitoring via a script within a batch file, $n = 1$, so this must be a matrix of a single row. If lags were included at the training step, then these observations will also have lagged features.
labelVector	an $n * 1$ integer vector of class memberships
trainingSummary	the TrainingSpecs list returned by the mspTrain() function. This list contains –for each class– the SPE and T2 thresholds, as well the projection matrix.
...	Lazy dots for additional internal arguments

Details

This function is designed to be ran at specific time intervals (every 10 seconds, 30 seconds, 1 minute, 5 minutes, 10 minutes, etc), from a batch file hosted script which calls this function and mspWarning(). This function takes in the specific observations to monitor and their class memberships (if any) and returns an xts matrix of these observations column concatenated with their monitoring statistic values, flag statuses, and an empty alarm column. Users should then append these rows onto a previously existing daily observations matrix. The mspWarning() function will then take in the daily observation xts matrix with updated rows returned by this function and check the monitoring statistic flag indicators to see if an alarm status has been reached. For further details, see the mspWarning() function.

This function calls the faultDetect() function, and requires the training information returned by the mspTrain function. This function will return the xts matrix necessary for the mspWarning() function.

Value

An $n * (p + 5)$ xts matrix, where the last five columns are the monitoring statistics and corresponding fault flags, and an empty alarm column

Examples

```
data("normal_switch_xts")
# The state values are recorded in the first column.
n <- nrow(normal_switch_xts)
nTrainObs <- floor(0.4 * n)

# Calculate the training summary, but save five observations for monitoring.
trainResults_ls <- mspTrain(data = normal_switch_xts[1:(n - 5), -1],
```

```

labelVector = normal_switch_xts[1:(n - 5), 1],
trainObs = nTrainObs,
lagsIncluded = 1)

# While training, we included 1 lag (the default), so we will also lag the
# observations we will test.
testObs <- normal_switch_xts[(n - 6):n, -1]
testObs <- stats::lag(testObs, 0:1)
testObs <- testObs[-1, ]
testObs <- cbind(normal_switch_xts[(n - 5):n, 1], testObs)

mspMonitor(observations = testObs[, -1],
            labelVector = testObs[, 1],
            trainingSummary = trainResults_ls$TrainingSpecs)

```

mspTrain

*Multi-State Adaptive-Dynamic Process Training***Description**

This function performs Multi-State Adaptive-Dynamic PCA on a data set with time-stamped observations.

Usage

```
mspTrain(data, labelVector, trainObs, updateFreq = ceiling(0.5 * trainObs),
        Dynamic = TRUE, lagsIncluded = 1, faultsToTriggerAlarm = 3, ...)
```

Arguments

data	An xts data matrix
labelVector	Class label vector (as logical or finite numeric)
trainObs	The number of observations upon which to train the algorithm
updateFreq	The algorithm update frequency (defaulting to half as many observations as the training frequency)
Dynamic	Should the PCA algorithm include lagged variables? Defaults to TRUE
lagsIncluded	If Dynamic = TRUE, how many lags should be included? Defaults to 1.
faultsToTriggerAlarm	The number of sequential faults needed to trigger an alarm
...	Lazy dots for additional internal arguments

Details

This function is designed to identify and sort out sequences of observations which fall outside normal operating conditions. We assume that the process data are time-dependent in both seasonal and non-stationary effects (which necessitate the Adaptive and Dynamic components, respectively). We further assume that this data is drawn from a multivariate process under multiple mutually exclusive states, implying that the linear dimension reduction projection matrices may be different for each state. Therefore, in summary, this function lags the features to account for correlation between sequential observations, splits the data by classes, and re-estimates projection matrices on a rolling

window to account for seasonality. Further, this function uses non-parametric density estimation to calculate the $1 - \alpha$ quantiles of the SPE and Hotelling's T2 statistics from a set of training observations, then flags any observation in the testing data set with process monitoring statistics beyond these calculated critical values. Because of natural variability inherent in all real data, we do not remove observations simply because they have been flagged as outside normal operating conditions. This function records an alarm only for observations having three flags in a row, as set by the default argument value of "faultsToTriggerAlarm". These alarm-positive observations are then removed from the data set and held in a separate xts matrix for inspection.

Of note when considering performance: the example has 10080 rows on three features alternating between three states, and trains on 20 percent of the observations, while updating every 1008 (10 percent) observation. On a 2016 Macbook Pro with 16Gb of RAM, this example function call takes 15 second to run. Increasing the update frequency will decrease computation time, but may increase false alarm rates or decrease flagging accuracy. We recommend that you set the update frequency based on the natural and physical designs of your system. For example, if your system has a multi-state process which switches across one of four states every two hours, then test the update frequency at an eight or 12 hour level — enough observations to measure two to three full cycles of the switching process. For observations recorded every five minutes, try $\text{updateFreq} = (60 / 5) * 8 = 96$ or $(60 / 5) * 12 = 144$.

This user-facing function calls the processMonitor() function, and returns the training arguments necessary to call the mspMonitor() and mspWarning() functions.

Value

A list of the following components: FaultChecks = an xts data matrix containing the SPE monitoring statistic and corresponding logical flagging indicator, the Hotelling's T2 monitoring statistic and corresponding logical flagging indicator, and the Alarm indicator. Non_Alarmed_Obs = an xts data matrix of all the non-Alarmed observations. Alarms = and an xts data matrix of the features and specific alarms for Alarmed observations, where the alarm code is as follows: 0 = no alarm, 1 = Hotelling's T2 alarm, 2 = Squared Prediction Error alarm, and 3 = both alarms. TrainingSpecs = a list of k lists, one for each class, with each list containing the specific threshold object returned by the internal threshold() function for that class. See this function's help file for more details.

Examples

```
data("normal_switch_xts")
nTrainObs <- floor(0.4 * nrow(normal_switch_xts))
# The state values are recorded in the first column.

mspTrain(data = normal_switch_xts[, -1],
          labelVector = normal_switch_xts[, 1],
          trainObs = nTrainObs)
```

mspWarning

Process Alarms

Description

Trigger an alarm, if necessary, for incoming multivariate process observations.

Usage

```
mspWarning(mspMonitor_object, faultsToTriggerAlarm = 3)
```

Arguments

- `mspMonitor_object`
An xts matrix returned by the `mspMonitor()` function
- `faultsToTriggerAlarm`
Specifies how many sequential faults will cause an alarm to trigger. Defaults to 3.

Details

This function and the `mspMonitor()` function are designed to be ran via a script within a batch. The file flow is as follows: at each time interval, run the `mspMonitor()` function on the daily observation matrix to add a flag status to the most recent incoming observation in the matrix, and return this new xts matrix. Then, pass this updated daily observation matrix to the `mspWarning()` function, which will check if the process has recorded three or more sequential monitoring statistic flags in a row. Of note, since these functions are expected to be repeatedly ran in real time, this function will only check for an alarm within the last row of the xts matrix. To check multiple rows for an alarm state, please use the `mspTrain` function, which was designed to check multiple past observations.

This function requires an xts matrix returned by the `mspMonitor()` function.

Value

An xts matrix of the same dimensions as `mspMonitor_object`, with a recorded negative or positive and type-specific alarm status

Examples

```
data("normal_switch_xts")
# The state values are recorded in the first column.
n <- nrow(normal_switch_xts)
nTrainObs <- floor(0.4 * n)

# Calculate the training summary, but save five observations for monitoring.
trainResults_ls <- mspTrain(data = normal_switch_xts[1:(n - 5), -1],
                           labelVector = normal_switch_xts[1:(n - 5), 1],
                           trainObs = nTrainObs,
                           lagsIncluded = 1)

# While training, we included 1 lag (the default), so we will also lag the
# observations we will test.
testObs <- normal_switch_xts[(n - 6):n, -1]
testObs <- stats::lag(testObs, 0:1)
testObs <- testObs[-1, ]
testObs <- cbind(normal_switch_xts[(n - 5):n, 1], testObs)

# Run the monitoring function.
dataAndFlags <- mspMonitor(observations = testObs[, -1],
                           labelVector = testObs[, 1],
                           trainingSummary = trainResults_ls$TrainingSpecs)

# Alarm check the last row of the matrix returned by the mspMonitor function
mspWarning(dataAndFlags)
```

normal_switch_xts

*Process Data under Normal Conditions***Description**

Three-feature, three-state process data under normal operating conditions as example data for different included functions.

Usage

```
normal_switch_xts
```

Format

An xts data matrix with 10080 rows, corresponding to one week worth of data recorded at a 1-minute interval, and four columns as defined here:

state the state indicator for the multivariate system, with three levels

x $x(t) = t + \text{error}$

y $y(t) = t^2 - 3t + \text{error}$

z $z(t) = -t^3 + 3t^2 + \text{error}$

where t is a 10080-entry vector of autocorrelated and non-stationary hidden process realizations. The states alternate each hour and are defined as follows:

State1 As presented

State2 Rotated by (yaw = 0, pitch = 90, roll = 30) and scaled by $(1 * x, 0.5 * y, 2 * z)$.

State3 Rotated by (yaw = 90, pitch = 0, roll = -30) and scaled by $(0.25 * x, 0.1 * y, 0.75 * z)$.

See the vignette for more details.

Source

Simulated in R.

pca

*PCA for Data Scatter Matrix***Description**

Calculate the principal components analysis for a data matrix, and also find the squared prediction error (SPE) and Hotelling's T2 test statistic values for each observation in this data matrix.

Usage

```
pca(data, var.amnt = 0.95, ...)
```

Arguments

<code>data</code>	A centred-and-scaled data matrix or xts matrix
<code>var.amnt</code>	How much energy should be preserved in the projection? Defaults to 0.95.
<code>...</code>	Lazy dots for additional internal arguments

Details

This function takes in a training data matrix, without the label column, and the energy preservation proportion, which defaults to 95 percent per Kazor et al (2016). This proportion is the sum of the q largest eigenvalues divided by the sum of all p eigenvalues, where q is the number of columns of the $p * q$ projection matrix P . This function then returns the projection matrix P , a diagonal matrix of the reciprocal eigenvalues (LambdaInv), a vector of the SPE test statistic values corresponding to the rows of the data matrix, and a T_2 test statistic vector similar to the SPE vector.

This internal function is called by `faultFilter()`.

Value

A list of class "pca" with the following: `projectionMatrix` = the q eigenvectors corresponding to the q largest eigenvalues as a $p * q$ projection matrix. `LambdaInv` = the diagonal matrix of inverse eigenvalues. `SPE` = the vector of SPE test statistic values for each of the n observations contained in the data matrix. `T2` = the vector of Hotelling's T_2 test statistic for each of the same n observations.

Examples

```
data("normal_switch_xts")
scaledData <- scale(normal_switch_xts[, -1])
pca(scaledData, var.amnt = 0.9)
```

processMonitor

Adaptive Process Training

Description

Apply Adaptive-Dynamic PCA to state-specific data matrices.

Usage

```
processMonitor(data, trainObs, updateFreq = ceiling(0.2 * trainObs),
  faultsToTriggerAlarm = 3, ...)
```

Arguments

<code>data</code>	An xts data matrix
<code>trainObs</code>	How many train observations will be used
<code>updateFreq</code>	How many non-flagged rows to collect before the function updates
<code>faultsToTriggerAlarm</code>	the number of sequential faults needed to trigger an alarm
<code>...</code>	Lazy dots for additional internal arguments

Details

This function is the class-specific implementation of the Adaptive- Dynamic PCA described in the details of the mspTrain function. See that function's help file for further details.

This internal function is called by mspTrain(). This function calls the faultFilter() function.

Value

A list of the following components: FaultChecks = a class specific xts data matrix containing the SPE monitoring statistic and corresponding logical flagging indicator, the Hotelling's T2 monitoring statistic and corresponding logical flagging indicator, and the Alarm indicator. Non_Alarmed_Obs = a class specific xts data matrix of all the observations with alarm states equal to 0. Alarms = a class-specific xts data matrix of the features and specific alarms for Alarmed observations, where the alarm code is as follows: 0 = no alarm, 1 = Hotelling's T2 alarm, 2 = Squared Prediction Error alarm, and 3 = both alarms. trainSpecs = the threshold object returned by the internal threshold() function. See this function's help file for more details.

Examples

```
data("normal_switch_xts")
# Select the data under state 1
data <- normal_switch_xts[normal_switch_xts[,1] == 1]
nTrainObs <- floor(0.4 * nrow(data))

# Remove the now unnecessary state column
featureCols <- data[,-1]

processMonitor(data = featureCols, trainObs = nTrainObs)
```

rotate3D

Three-Dimensional Rotation Matrix

Description

Render a 3-Dimensional projection matrix given positive or negative degree changes in yaw, pitch, and / or roll.

Usage

```
rotate3D(yaw, pitch, roll)
```

Arguments

yaw	z-axis change in degrees; look left (+) or right (-). Consider this a rotation on the x-y plane.
pitch	y-axis change in degrees; look up (-) or down (+). Consider this a rotation on the x-z plane.
roll	x-axis change in degrees; this change appears as if you touch head to shoulders: right roll (+) and left roll (-).

threshold	<i>Non-parametric Threshold Estimation</i>
-----------	--

Description

Calculate the non-parametric critical value estimates for the SPE and T2 monitoring test statistics.

Usage

```
threshold(pca_object, alpha = 0.001, ...)
```

Arguments

pca_object	A list with class "pca" from the internal pca() function
alpha	The upper 1 - alpha quantile of the SPE and T2 densities from the training data passed to this function. Defaults to 0.001.
...	Lazy dots for additional internal arguments

Details

This function takes in a pca object returned by the pca() function and a threshold level defaulting to 0.1 quantile is set this low to reduce false alarms, as described in Kazor et al (2016). The function then returns a calculated SPE threshold corresponding to the 1 - alpha critical value, a similar T2 threshold, and the projection and Lambda Inverse (1 / eigenvalues) matrices passed through from the pca() function call.

This internal function is called by faultFilter().

Value

A list with classes "threshold" and "pca" containing: SPE_threshold = the 1 - alpha quantile of the SPE density. T2_threshold = the 1 - alpha quantile of the Hotelling's T2 density. projectionMatrix = a projection matrix from the data feature space to the feature subspace which preserves some specified proportion of the energy of the data scatter matrix. This specified energy proportion is user specified through the var.amnt argument in the pca() function. LambdaInv = a diagonal matrix of the reciprocal eigenvalues of the data scatter matrix.

Examples

```
data("normal_switch_xts")
scaledData <- scale(normal_switch_xts[, -1])
pca_obj <- pca(scaledData, var.amnt = 0.9)
threshold(pca_object = pca_obj, alpha = 0.05)
```

Index

*Topic **datasets**

- fault1A_xts, [2](#)
- fault1B_xts, [3](#)
- fault2A_xts, [3](#)
- fault2B_xts, [4](#)
- fault3A_xts, [5](#)
- fault3B_xts, [6](#)
- normal_switch_xts, [13](#)

- fault1A_xts, [2](#)
- fault1B_xts, [3](#)
- fault2A_xts, [3](#)
- fault2B_xts, [4](#)
- fault3A_xts, [5](#)
- fault3B_xts, [6](#)
- faultDetect, [6](#)
- faultFilter, [7](#)

- mSPMonitor, [9](#)
- mSPTrain, [10](#)
- mSPWarning, [11](#)

- normal_switch_xts, [13](#)

- pca, [13](#)
- processMonitor, [14](#)

- rotate3D, [15](#)
- rotateScale3D, [16](#)

- threshold, [17](#)